

Routing-Level Signatures of Cognitive Efficiency: Only Inter-Subject Routing Consistency Survives FDR, with a Pairwise-Correlation Caveat

Xiaoyan Li*, Cuicui Jiang[†], Jiaoping Chen[‡], Yujia Du[§], Jiaxuan Wei[¶], Xingyue Liu[¶], and Rumei Yang^{||}

*Department of Computational Mathematics, Science and Engineering, Michigan State University, East Lansing, MI, USA

[†]College of Computer Science, Inner Mongolia University, Hohhot, China

[‡]Merrick School of Business, University of Baltimore, Baltimore, MD, USA

[§]Suzhou MetaCortex Intelligence Technology Co., Ltd., Suzhou, China

[¶]School of Artificial Intelligence, Taiyuan University of Technology, Jinzhong, China

^{||}School of Nursing, Nanjing Medical University, Nanjing, China

lixiaoy5@msu.edu, jiangcuicui@mail.imu.edu.cn, jchen@ubalt.edu, saroyi@163.com,
2023002323@link.tyut.edu.cn, 2023002320@link.tyut.edu.cn, rumeiyang@njmu.edu.cn

Corresponding author: Rumei Yang (rumeiyang@njmu.edu.cn)

Abstract—We jointly test four pre-specified routing-level candidate signatures of cognitive efficiency (activation-distribution sparsity, canonical task-network engagement, inter-subject routing consistency, and modularity) in 100 Human Connectome Project Young Adult participants across all seven task-fMRI paradigms, with efficiency groups defined by working-memory behavior and Benjamini–Hochberg (BH) FDR control. Only inter-subject routing consistency survives FDR: under a subject-level leave-one-out (LOO) statistic, high-efficiency participants route more like one another ($d = +0.52$, 95% CI $[+0.12, +0.92]$, $p = 0.011$, $p_{\text{FDR}} = 0.033$; label permutation $p = 0.018$; $\text{BF}_{10} = 4.0$), and consistency correlates with the continuous efficiency score ($r = 0.36$, $p = 0.0002$). Sparsity ($d = -0.26$), task-network engagement (2/7 canonical matches, at chance) and modularity ($d = +0.19$) are not supported. On the same consistency data, the conventional within-group pairwise-correlation t -test reports $p = 1.0 \times 10^{-30}$, overstating significance by about 28 orders of magnitude because it treats 2450 correlated pairs as independent; we show the inflation directly and give a drop-in subject-level replacement. A cross-domain comparison with three Mixture-of-Experts (MoE) language models finds cortex about $3.3\times$ more concentrated than MoE expert use pooled over layers, but single MoE layers can be far more concentrated, so that comparison depends on the level of pooling. Routing consistency is a candidate marker of efficient cognition that now needs independent replication.

Index Terms—neural efficiency, multi-task fMRI, inter-subject routing consistency, pairwise-correlation inference, FDR

I. INTRODUCTION

Cognitive efficiency varies substantially between individuals and declines with age [1], [2], with pronounced deficits in Alzheimer’s disease [3]–[5] and schizophrenia [6]. Translating this heterogeneity into individual-subject prediction is a central goal of “precision neuroimaging” [7], [8], and task fMRI is an attractive substrate because it probes how information is routed across functional networks, the system whose selective, metabolically bounded engagement is thought to underlie efficient cognition [2], [9].

The neural efficiency hypothesis proposes that more capable brains accomplish tasks with fewer neural resources [2], [10], and sparse coding theory formalizes the idea that information is carried by selective activation of small neuronal subsets [11], [12]. A routing-based account must also say *which* circuits are engaged for which task, *how consistent* that routing is across individuals, and *how modular* the resulting organization is, each with its own clinical reading, for example loss of network segregation in Alzheimer’s disease [3], [4] or diffuse recruitment during working memory in schizophrenia [6]. Individual-differences studies have usually examined one of these metrics at a time: modularity as a marker of plasticity [13], [14], connectivity fingerprinting [7], [8], or activation extent [10]. To our knowledge the four properties (sparser, more task-selective, more consistent, more modular) have not been tested jointly in one sample with a behavior-derived efficiency label and multiple-comparison control.

We also compare cortex with three Mixture-of-Experts (MoE) language models [15], [16], which obtain computational efficiency by sparse expert routing [17]. Earlier brain–network comparisons concern representational similarity [18]–[20] rather than routing.

The paper makes three contributions. (1) In 100 HCP participants, with the HCP-MMP1.0 parcellation, task-canonical GLM contrasts, Louvain community detection and BH-FDR-controlled joint testing, one of the four candidates, inter-subject routing consistency, survives FDR, and permutation, rank, continuous-score and Bayesian analyses agree; sparsity, canonical task-network engagement and modularity are not supported. (2) We quantify an inference pitfall that the consistency analysis invites: the within-group pairwise-correlation t -test treats $n(n-1)/2$ non-independent pairs as independent and, on our data, overstates significance by about 28 orders of magnitude relative to subject-level and permutation inference (Fig. 2); we give the subject-level statistic as a drop-

in replacement. (3) Cortex is more concentrated than MoE expert use pooled over layers, but single MoE layers can be more concentrated than cortex, so cross-domain efficiency comparisons must state the level of pooling.

II. METHODS

A. Participants, Data, and Preprocessing

We used HCP Young Adult task fMRI [21] from 100 healthy adults (53 female; open-access age bands 22–25: 17, 26–30: 42, 31–35: 40, > 35: 1). The sample is the first 100 participants, in ascending subject ID, of a list of HCP-YA participants with complete data for all seven tasks in both phase-encoding directions (LR and RL); no further selection was applied, and no motion threshold was used for selection. Head motion was nonetheless low: mean relative RMS displacement was 0.089 mm (SD 0.030), no run exceeded 0.5 mm (maximum 0.36 mm), and motion did not differ between efficiency groups ($p = 0.64$). We used the MSMAll-aligned, minimally preprocessed CIFTI time series (tfMRI*_Atlas_MSMAll.dtseries.nii) [22], [23], without ICA-FIX denoising. HCP-YA contains twins and siblings; family identifiers are restricted data that we did not use, so family structure is not modelled (Sec. IV-D).

B. Parcellation and Activation Extraction

Each run’s grayordinate time series was averaged within the 360 cortical parcels of HCP-MMP1.0 [24]. For each run we fitted a GLM to the parcel time series containing only the HCP block regressors of that task, convolved with the SPM canonical double-gamma haemodynamic response (TR = 0.72 s), a discrete-cosine high-pass basis (128 s) and an intercept. Contrasts followed the HCP task definitions: 2-back vs. 0-back (WM); all five movements vs. fixation (MOTOR); story vs. math (LANGUAGE); mental vs. random motion (SOCIAL); relational vs. match (RELATIONAL); fearful faces vs. shapes (EMOTION); win vs. loss (GAMBLING). The LR and RL contrast vectors were averaged and z -scored across parcels, giving a $100 \times 7 \times 360$ activation array; all 700 subject $\times$ task cells were extracted, and design condition numbers were at most 1.7. This first-level model corrects the submitted version, whose GLM also contained event-level regressors overlapping the task blocks (for WM, a rank-deficient design), matched regressors by name substring, and contrasted MOTOR against the cue; all results below come from the corrected model. Each parcel was assigned to one of the 12 Cole-Anticevic networks [25] by modal vote, and the 12 were collapsed to seven Yeo-like networks [26]: {Visual1, Visual2} $\rightarrow$ VIS; {Somatomotor, Auditory} $\rightarrow$ SMN; {Dorsal-Attention} $\rightarrow$ DAN; {Cingulo-Opercular, Language, Ventral-Multimodal} $\rightarrow$ VAN; {Orbito-Affective} $\rightarrow$ LIM; {Frontoparietal, Posterior-Multimodal} $\rightarrow$ FPN; {Default} $\rightarrow$ DMN (60, 54, 23, 83, 6, 57 and 77 parcels).

C. Efficiency Score

Participants were split at the median of a working-memory accuracy-to-reaction-time ratio, $E_i = \text{Acc}_i / (\text{RT}_i / \text{RT})$, with Acc_i the mean 0-back/2-back accuracy and RT_i the mean median response time, into high- and low-efficiency groups ($n = 50$ each). The score is independent of the neuroimaging measures, being computed from behavior alone, but not of the task battery: Working Memory is also one of the seven fMRI paradigms. Its run-to-run reliability (LR vs. RL) was $r = 0.80$ (Spearman–Brown 0.89), and 84% of participants fall on the same side of the median in both runs. Because a median split discards information, every primary measure is also related to the continuous E_i (Sec. III-E).

D. Hypotheses, Primary Outcomes, and Inference

Each hypothesis has one primary outcome and one test, computed with one value per participant. **H1** (efficient brains are sparser): composite sparsity averaged over the seven tasks, Student t . **H2** (task-selective engagement): the descriptive canonical-match count below and, added for this version, a group test of each participant’s mean entropy selectivity (Welch t). **H3** (more consistent routing): subject-level LOO correlation, Welch t . **H4** (more modular): Newman’s Q , Student t . The pre-specified primary BH-FDR family is {H1, H3, H4}; because H2 had no group test in the original plan, it enters only a four-test sensitivity family. Secondary families (the four H1 metrics, the three H4 metrics, the seven per-task H1 tests and the six continuous-score correlations) are corrected within themselves. For each primary test we report Cohen’s d with a 95% CI, a two one-sided test (TOST) of equivalence to $|d| < 0.5$ [27], and a JZS Bayes factor BF_{10} (Cauchy prior scale $\sqrt{2}/2$) [28]. The submitted version tested H1 on all 700 subject $\times$ task values as if independent; the subject-level test used here respects the independence requirement that motivates Sec. II-F.

E. Activation-Distribution Sparsity

Three indices describe how concentrated a task map is across parcels: the Gini coefficient of $u_i = |a_i|$, $G = \sum_{i,j} |u_i - u_j| / (2n \sum_i u_i)$; an entropy index $S = 1 - H / \ln B$ from a $B = 50$ -bin histogram of the z -scored values; and the share of $\sum_i u_i$ carried by the top 10% of parcels. The composite is their mean. Because maps are z -scored within subject and task, these indices measure the spatial concentration of the activation distribution, not activation extent, amplitude or metabolic cost; we call the construct activation-distribution sparsity and do not equate it with neural resource consumption. z -scoring also makes L_0 sparsity and supra-threshold region counts constant, so they are not used.

F. Routing Consistency and the Subject-Level Statistic

Each participant’s routing profile $\mathbf{x}_i$ is the flattened 7×7 task-by-network matrix of network-mean activations. The primary H3 statistic is

$$r_i^{\text{LOO}} = \text{corr} \left(\mathbf{x}_i, \frac{1}{|g_i|-1} \sum_{j \in g_i, j \neq i} \mathbf{x}_j \right), \quad (1)$$

where g_i is i 's group; the 100 subject-level LOO statistics enter a Welch t -test. In pseudocode: `for i in group: r[i] = corr(x[i], mean(x[group \ i])); welch_t(r[high], r[low])`. The LOO templates of members of the same group overlap, so the r_i are not strictly independent; a label-permutation test (10000 permutations, statistic = difference of within-group mean pairwise r), which makes no independence assumption, is the confirmatory check. For comparison we also run the analysis a practitioner would reach for, and which our own earlier draft used: all 1225 within-group pairwise correlations per group pooled into an independent-samples t -test.

G. Task Selectivity and Modularity

Task selectivity is $S_t = 1 - H_t / \log_2 K$, with H_t the entropy of the absolute network-mean activations normalized to proportions over K networks ($K = 7$, or $K = 12$ at native resolution). Canonical expectations were fixed before analysis at Yeo-7 (WM, RELATIONAL: FPN or DAN; MOTOR: SMN; LANGUAGE: FPN or VAN; SOCIAL: DMN; EMOTION: LIM; GAMBLING: FPN); at 12 networks each expectation is carried over to the Cole-Anticevic networks that collapse into it, so no new choice is made. A task matches if its highest signed network mean is an expected network; chance is estimated by permuting network labels.

For H4, each participant's 360×360 matrix A holds Pearson correlations between parcels' seven-task activation vectors. This is task-activation-profile similarity with seven observations per edge, not time-series functional connectivity. Louvain [29] (resolution 1, seed 0) is run on $|A|$: two parcels whose task profiles are strongly anti-correlated are coupled by the same task factor with opposite sign, so we treat them as related rather than unrelated; a positive-edge-only graph is a sensitivity check. From the partition c we compute Newman's $Q = \frac{1}{2m} \sum_{ij} [|A_{ij}| - k_i k_j / 2m] \delta(c_i, c_j)$, segregation $(\bar{w}_{\text{within}} - \bar{w}_{\text{between}}) / (\bar{w}_{\text{within}} + \bar{w}_{\text{between}})$ with $\bar{w}$ the mean $|A_{ij}|$ within and between communities, and the mean participation coefficient $P_i = 1 - \sum_c (k_{ic} / k_i)^2$ [30]. Partition stability is assessed over 20 Louvain seeds.

H. Mixture-of-Experts Models

We analyzed Switch-base-8 (8 experts, Top-1) [16], Qwen1.5-MoE-A2.7B (60 routed experts, Top-4) [31] and DeepSeek-MoE-16B-base (64 routed experts, Top-6) [32]. Text came from RACE [33], SNLI [34], SST-2 [35] and GSM8K [36], sorted into seven categories each paired with one fMRI task (Table I); 618 samples were collected and up to 50 per category (318 sequences, ≤ 128 tokens) were run through each model. Routing was read from the router modules by forward hooks. The extractor of the submitted version also matched non-router modules whose names contain "gate" and counted their outputs as expert choices; all MoE numbers here come from a corrected extractor that accepts only per-token distributions over the E experts or the router's selected expert indices. For each model we compute the Gini coefficient of expert use pooled over layers and tokens, the same per

TABLE I: Correspondence between fMRI tasks and MoE text categories. The pairing is an analogy, not a matched manipulation; RISK-GAMBLING is the weakest.

fMRI task	Text category	Source (selection)
Working Memory	Mathematical	GSM8K word problems
Motor	Procedural	RACE passages (instructional)
Language	Narrative	RACE passages (story)
Social	Social	SNLI premises (social keywords)
Relational	Scientific	RACE passages (expository)
Emotion	Emotional	SST-2 sentences
Gambling	Risk decision	SNLI premises (risk keywords)

TABLE II: Activation-distribution sparsity by efficiency group, one value per participant (mean over seven tasks). Negative d : the high-efficiency group is less sparse. p_{FDR} is within the four-metric H1 family.

Metric	High	Low	d	p	p_{FDR}
Gini coefficient	0.445	0.445	-0.05	0.80	0.80
Entropy sparsity	0.157	0.163	-0.35	0.082	0.33
Concentration ratio	0.288	0.290	-0.21	0.30	0.40
Composite	0.297	0.299	-0.26	0.19	0.39

layer (Switch: encoder layers only, since the decoder sees one token per sequence), the effective number of experts e^H , and each expert's dominant-category share (uniform = $1/7$). At the token level routing is sparse by construction: each token activates k of E experts.

III. RESULTS

A. H1: Activation-Distribution Sparsity

Sparsity did not differ between groups (Table II; composite $d = -0.26$, 95% CI $[-0.66, +0.13]$, $p = 0.19$). Every index pointed, if anything, toward *less* sparse activation in the high-efficiency group, the direction opposite to the neural-efficiency prediction, and none of the seven per-task tests crossed FDR (smallest $p_{\text{FDR}} = 0.49$). H1 is not supported.

B. H2: Task-Selective Network Engagement

Parcel-level maps show the expected anatomy (Fig. 1): WM and RELATIONAL engage lateral frontoparietal cortex, MOTOR the central sulcus and EMOTION ventral visual cortex. At the network level, mean selectivity at Yeo-7 was $\bar{S}_t = 0.106$, close to the uniform baseline, and the most engaged network matched the canonical prediction in 2/7 tasks (WM and RELATIONAL, both FPN). At the native 12 networks the count was also 2/7 (the same two tasks, Frontoparietal), against a chance expectation of 1.4 (permutation $p = 0.44$); MOTOR peaked in Cingulo-Opercular with Somatomotor second, and LANGUAGE in Default with Language second. The signed network mean spreads focal peaks over large networks, so the match count is a coarse test, but it is not rescued by the finer partition. Selectivity did not differ between efficiency groups (Yeo-7: $d = +0.14$, $p = 0.50$; 12 networks: $d = +0.40$, $p = 0.050$). H2 is not supported.




Fig. 1: Group-averaged task-evoked activation for the seven HCP tasks ($N = 100$) on the HCP-MMP1.0 atlas [24] in MNI152 space (dorsal glass-brain views, one shared diverging color scale). The most engaged network matches the canonical prediction in 2/7 tasks at both Yeo-7 and 12-network resolution (WM and RELATIONAL, frontoparietal).

C. H3: Inter-Subject Routing Consistency

Primary. Subject-level LOO correlations were $r = 0.680$ (high) and 0.626 (low): $d = +0.52$, 95% CI $[+0.12, +0.92]$, Welch $t = 2.59$, $p = 0.011$; Mann-Whitney $p = 0.013$; bootstrap CI on the difference $[+0.014, +0.092]$. The label-permutation test, which needs no independence assumption, agrees ($\Delta r = +0.069$, null SD 0.029 , $p = 0.018$). H3 survives the primary BH-FDR family ($p_{\text{FDR}} = 0.033$) and the four-test sensitivity family ($p_{\text{FDR}} = 0.044$), and the Bayes factor gives moderate evidence for a difference ($\text{BF}_{10} = 4.0$). Consistency is lower in participants who move more ($r = -0.28$ with mean relative RMS), but motion does not differ between groups, and with motion as a covariate the group effect is unchanged ($p = 0.013$). The CI is wide: effects from $d \approx 0.1$ to 0.9 are compatible with the data, and resolving the lower end would need about 1100 participants per group, whereas replicating $d = 0.52$ at 80% power needs about 60.

The pairwise-correlation caveat. Pooling the 1225 within-group pairwise correlations per group into an independent-samples t -test gives, on the same data, $t = 11.68$ and $p = 1.0 \times 10^{-30}$ with an apparent $d = +0.47$. The test inflates the effective sample size, from 100 participants to about 2450 pairs, and overstates significance by about 28 orders of magnitude. Fig. 2 shows why: the null distribution implied by the naive test is about $4.9\times$ narrower than the permutation null. Here the conclusion survives correct inference, but only just; on the submitted first-level model the same naive test gave $p = 10^{-17}$ while subject-level inference gave $p = 0.049$ and permutation $p = 0.055$, so the naive test would report an overwhelming effect in either case, regardless of whether a correct test finds one.

D. H4: Network Modularity

Louvain found 3–6 communities per participant (median 4). Modularity did not differ between groups (Q : $d = +0.19$, $p = 0.34$; Fig. 3), nor did segregation ($d = +0.35$, $p = 0.085$) or participation ($d = -0.14$, $p = 0.50$; all within-family $p_{\text{FDR}} \geq 0.25$). Individual partitions varied across 20 Louvain seeds (median ARI 0.54 , NMI 0.54), but Q was stable (within-participant SD 0.002 against a between-participant SD of




Fig. 2: Null distributions for the H3 group difference on one axis. Top: normal null implied by the naive pairwise t -test. Middle: null of the subject-level LOO Welch test. Bottom: label-permutation null (10 000 permutations). Dashed line: observed difference. The naive null is far too narrow because the 2450 pairs share 100 participants.

0.007), and seed-averaged tests (Q : $d = +0.25$, $p = 0.21$) and a positive-edge-only graph ($d = +0.33$, $p = 0.11$) reproduce the null. H4 is not supported.

E. Sensitivity Analyses

With the continuous efficiency score in place of the median split (Table III), routing consistency, computed against the whole-sample LOO template, correlated positively with efficiency ($r = 0.36$, 95% CI $[0.18, 0.52]$, $p = 0.0002$, $p_{\text{FDR}} = 0.001$; permutation of E_i : $p = 0.0002$; partial $r = 0.34$ controlling for motion). Two secondary associations also crossed FDR within this family: composite sparsity ($r = -0.26$, $p_{\text{FDR}} = 0.026$) and participation ($r = -0.24$, $p_{\text{FDR}} = 0.037$), both negative, i.e. more efficient participants




Fig. 3: Network-organization metrics by efficiency group (Louvain, seed 0): modularity Q , segregation and mean participation coefficient (half-violin, box plot and one dot per participant). Cohen’s d is printed in each panel; none of the three metrics crosses FDR.

TABLE III: Summary of the brain-side hypotheses. d : high minus low efficiency, with 95% CI. p_{FDR} : primary family $\{H1, H3, H4\}$; $H2$ is in the four-test sensitivity family. TOST: p for equivalence to $|d| < 0.5$. r_E : Pearson correlation with the continuous efficiency score (p_{FDR} across six measures; for $H3$ against the whole-sample LOO template). $H3$ uses the subject-level LOO statistic; the naive pairwise test on the same data gives $p = 1.0 \times 10^{-30}$.

	Primary outcome	Verdict	d [95% CI]	p	p_{FDR}	TOST p	BF_{10}	r_E (p_{FDR})
H1	Composite sparsity (subject mean)	Not supported	-0.26 $[-0.66, +0.13]$	0.19	0.29	0.12	0.45	-0.26 (0.026)
H2	Task selectivity $\bar{S}_t$ (Yeo-7)	Not supported	$+0.14$ $[-0.26, +0.53]$	0.50	0.50	0.036	0.26	-0.02 (0.82)
H3	LOO routing consistency	Supported	$+0.52$ $[+0.12, +0.92]$	0.011	0.033	0.54	4.0	$+0.36$ (0.001)
H4	Modularity Q	Not supported	$+0.19$ $[-0.20, +0.59]$	0.34	0.34	0.064	0.32	$+0.18$ (0.10)

distribute activation slightly more broadly and have slightly less cross-community integration; neither was detected by the pre-specified group tests. Modularity ($r = 0.18$) and selectivity ($r = -0.02$) did not.

F. MoE Routing and the Cross-Domain Comparison

Pooled over layers and tokens, all three models use their experts almost uniformly (Table IV): aggregate Gini 0.122, 0.038 and 0.247 (mean 0.135), against a cortical Gini of 0.445 over 360 parcels, a ratio of about $3.3\times$. The difference is not an artifact of the number of units: cortical Gini on $k = 8, 60$ or 64 randomly chosen parcels, matching each model’s expert count, is 0.39, 0.44 and 0.44 (Fig. 4). No expert specialized for a category; the dominant-category share stayed near the uniform $1/7$ (0.24–0.25 in all three models).

Pooling, however, hides concentration. Within a single MoE layer expert use is much more concentrated than pooled use (mean per-layer Gini 0.29, 0.17 and 0.76), because different layers favor different experts and the preferences cancel when layers are pooled; DeepSeek’s single layers are more concentrated than cortex (Fig. 4). At the token level routing is sparse by design, with k/E between 6.7% and 12.5% of experts active per token. Which system is “sparser” therefore depends on the level of comparison, and the $3.3\times$ figure applies only to pooled expert use. We report no inferential test for the cortex–MoE difference: there are only three models, and the submitted

TABLE IV: MoE routing with the corrected router extraction. Pooled: Gini of expert use pooled over layers and tokens; per layer: mean [min, max] over layers; $E_{\text{eff}} = e^H$ of pooled use; Dom.: mean dominant-category share (uniform $1/7 \approx 0.143$).

Model	k/E	Pooled	Per layer	E_{eff}	Dom.
Switch-8	1/8	0.122	0.29 [0.16, 0.45]	7.8/8	0.248
Qwen-MoE	4/60	0.038	0.17 [0.12, 0.24]	59.9/60	0.247
DeepSeek-16B	6/64	0.247	0.76 [0.33, 0.86]	57.9/64	0.244
Cortex (360 parcels)	—	0.445	—	—	—

Cohen’s $d = 4.69$ treated 700 subject $\times$ task brain values and three model values as a two-sample comparison, which has no meaningful sampling unit.

The submitted structural-similarity correlation ($r = +0.25$, $p = 0.089$) correlated, element by element, the flattened z -scored 7×7 brain task-by-network matrix with the first 49 entries of the 7×8 Switch category-by-expert matrix. Because expert indices are arbitrary, that number is not interpretable. We replace it with a permutation-invariant comparison of the 7×7 task-similarity matrix of cortex with each model’s 7×7 category-similarity matrix (Mantel test [37] over all 5040 label orderings): Spearman $r = -0.18, -0.41$ and -0.28 (exact $p \geq 0.77$). Cortex and MoE routing share no detectable category structure.




Fig. 4: Cortex versus MoE concentration at matched unit counts. Squares: cortical Gini on k randomly chosen parcels (mean and 95% range over participants $\times$ tasks), with k equal to each model’s number of experts; dotted line: all 360 parcels. Diamonds: MoE expert use pooled over layers; open circles: single MoE layers.

IV. DISCUSSION

A. What Distinguishes Efficient Routing

Of four candidate signatures, only one survives joint FDR control: high-efficiency participants route task information across networks more like one another. The effect is moderate, is carried by both the group test and the continuous score, and holds under rank and permutation inference, which is what distinguishes it from the pairwise-correlation evidence often offered for such claims. Sparsity does not separate the groups, and the continuous analysis points away from the classical account: more efficient participants distribute activation slightly more broadly, in line with the reading of neural efficiency as targeted rather than minimal activation [2], [10]. Canonical task-network engagement is at chance even at 12 networks, and modularity is null across seeds and edge treatments, consistent with reports that modularity–cognition associations depend on partitioning and edge weighting [13], [14]. Consistency, not concentration, is where efficiency shows up in these data.

B. Implications for Routing-Level Biomarkers

Routing consistency is a candidate marker, not a validated one. Its CI spans small to large effects, it comes from one cohort, and our efficiency score uses one of the seven tasks. A replication with about 60 participants per group would test the observed effect; ruling out a small one needs far more. Any such replication, and any claim of “more consistent routing” in preclinical Alzheimer’s, schizophrenia [3], [4], [6] or rehabilitation cohorts, should report a subject-level or permutation test: on our data the naive pairwise test overstates significance by about 28 orders of magnitude, and on the submitted first-level model it reported $p = 10^{-17}$ for an effect that subject-level inference placed at $p \approx 0.05$.

C. What the MoE Comparison Does and Does Not Show

Pooled over layers, MoE expert use is near-uniform and category-blind, as expected from load-balancing objectives [16], but single layers can be more concentrated than cortex. Moreover, our efficiency signal lies in the consistency of routing across individuals, which has no direct analogue in a single trained model; the cross-domain comparison therefore concerns concentration, not the property that distinguishes efficient brains here.

D. Limitations

(i) Family structure is not modelled; HCP-YA siblings and twins may make participants non-independent, which could inflate the H3 evidence. Consistency also depends on head motion, which we controlled for statistically but not by design. (ii) The data were not ICA-FIX denoised, and no motion threshold was applied, although motion was low and balanced across groups. (iii) The efficiency score comes from one session and one of the seven tasks; its run-to-run reliability is 0.80. (iv) H4 uses task-profile similarity with seven observations per edge, not time-series FC, and results may differ with Leiden refinement [38]. (v) The task–category mapping (Table I) is approximate, and the MoE results apply to the three examined architectures. (vi) The sample is a narrow healthy young-adult cohort. The pairwise-correlation pitfall is a property of the statistical procedure and does not depend on these.

V. CONCLUSION

In 100 HCP young adults, joint FDR-controlled testing of four routing-level signatures of cognitive efficiency supports one: high-efficiency participants route task information more consistently with one another ($d = +0.52$, $p_{\text{FDR}} = 0.033$), a result that holds under permutation, rank, continuous-score and Bayesian analysis. Sparsity, canonical task-network engagement and modularity are not supported. The conventional pairwise-correlation t -test would have reported the consistency effect at $p = 10^{-30}$; subject-level and permutation inference put it at $p \approx 0.01$ – 0.02 , and on a flawed first-level model the naive test was just as confident while correct inference was not. Cortex is more concentrated than pooled MoE expert use, but not than every single MoE layer. Routing consistency is a candidate marker of efficient cognition that now needs independent replication with subject-level inference.

ACKNOWLEDGMENT

Data were provided by the Human Connectome Project, WU-Minn Consortium (Principal Investigators: David Van Essen and Kamil Ugurbil; 1U54MH091657) funded by the 16 NIH Institutes and Centers that support the NIH Blueprint for Neuroscience Research, and by the McDonnell Center for Systems Neuroscience at Washington University. Claude (Anthropic) was used to assist with analysis code, manuscript revision and language editing; the authors verified all content and take full responsibility for it.

CODE AND DATA AVAILABILITY

Analysis code, group-level results and the MoE routing outputs are available at <https://github.com/xiaoyanLi629/routing-signatures-efficiency>. HCP data are available from the Human Connectome Project under its open-access terms and are not redistributed.

REFERENCES

- [1] M. Y. Chan, D. C. Park, N. K. Savalia, S. E. Petersen, and G. S. Wig, "Decreased segregation of brain systems across the healthy adult lifespan," *Proceedings of the National Academy of Sciences*, vol. 111, no. 46, pp. E4997–E5006, 2014.
- [2] A. C. Neubauer and A. Fink, "Intelligence and neural efficiency," *Neuroscience & Biobehavioral Reviews*, vol. 33, no. 7, pp. 1004–1023, 2009.
- [3] Z. Zhang, M. Y. Chan, L. Han, C. A. Carreno, E. Winter-Nelson, and G. S. Wig, "Dissociable effects of Alzheimer's disease-related cognitive dysfunction and aging on functional brain network segregation," *Journal of Neuroscience*, vol. 43, no. 46, pp. 7879–7892, 2023.
- [4] M. Ewers, Y. Luan, L. Frontzkowski, J. Neitzel, A. Rubinski, M. Dichgans, J. Hassenstab, B. A. Gordon, J. P. Chhatwal, J. Levin, and N. Franzmeier, "Segregation of functional networks is associated with cognitive resilience in Alzheimer's disease," *Brain*, vol. 144, no. 7, pp. 2176–2185, 2021.
- [5] J. Gonneaud, A. T. Baria, A. Pichet Binette, B. A. Gordon, J. P. Chhatwal, C. Cruchaga, M. Jucker, J. Levin, S. Salloway, M. Farlow, S. Gauthier, T. L. S. Benzinger, J. C. Morris, R. J. Bateman, J. C. S. Breitner, J. Poirier, E. Vachon-Presseau, and S. Villeneuve, "Accelerated functional brain aging in pre-clinical familial Alzheimer's disease," *Nature Communications*, vol. 12, no. 1, p. 5346, 2021.
- [6] M. A. Kim, E. Tura, S. G. Potkin, M. H. Fallon, D. S. Manoach, V. D. Calhoun, and J. A. Turner, "Working memory circuitry in schizophrenia shows widespread cortical inefficiency and compensation," *Schizophrenia Research*, vol. 117, no. 1, pp. 42–51, 2010.
- [7] E. S. Finn, X. Shen, D. Scheinost, M. D. Rosenberg, J. Huang, M. M. Chun, X. Papademetris, and R. T. Constable, "Functional connectome fingerprinting: identifying individuals using patterns of brain connectivity," *Nature Neuroscience*, vol. 18, no. 11, pp. 1664–1671, 2015.
- [8] X. Shen, E. S. Finn, D. Scheinost, M. D. Rosenberg, M. M. Chun, X. Papademetris, and R. T. Constable, "Using connectome-based predictive modeling to predict individual behavior from brain connectivity," *Nature Protocols*, vol. 12, no. 3, pp. 506–518, 2017.
- [9] D. Attwell and S. B. Laughlin, "An energy budget for signaling in the grey matter of the brain," *Journal of Cerebral Blood Flow & Metabolism*, vol. 21, no. 10, pp. 1133–1145, 2001.
- [10] R. J. Haier, B. V. Siegel Jr, K. H. Nuechterlein, E. Hazlett, J. C. Wu, J. Paek, H. L. Browning, and M. S. Buchsbaum, "Cortical glucose metabolic rate correlates of abstract reasoning and attention studied with positron emission tomography," *Intelligence*, vol. 12, no. 2, pp. 199–217, 1988.
- [11] B. A. Olshausen and D. J. Field, "Emergence of simple-cell receptive field properties by learning a sparse code for natural images," *Nature*, vol. 381, no. 6583, pp. 607–609, 1996.
- [12] —, "Sparse coding of sensory inputs," *Current Opinion in Neurobiology*, vol. 14, no. 4, pp. 481–487, 2004.
- [13] C. L. Gallen and M. D'Esposito, "Brain modularity: A biomarker of intervention-related plasticity," *Trends in Cognitive Sciences*, vol. 23, no. 4, pp. 293–304, 2019.
- [14] K. L. Arnamann, A. J.-W. Chen, T. Novakovic-Agopian, C. Gratton, E. M. Nomura, and M. D'Esposito, "Functional brain network modularity predicts response to cognitive training after brain injury," *Neurology*, vol. 84, no. 15, pp. 1568–1574, 2015.
- [15] N. Shazeer, A. Mirhoseini, K. Maziarz, A. Davis, Q. Le, G. Hinton, and J. Dean, "Outrageously large neural networks: The sparsely-gated mixture-of-experts layer," *arXiv preprint arXiv:1701.06538*, 2017.
- [16] W. Fedus, B. Zoph, and N. Shazeer, "Switch transformers: Scaling to trillion parameter models with simple and efficient sparsity," *Journal of Machine Learning Research*, vol. 23, no. 120, pp. 1–39, 2022.
- [17] R. Bommasani, D. A. Hudson, E. Adeli, R. Altman, S. Arber, S. von Arx, M. S. Bernstein, J. Bohg, A. Bosselut, E. Brunskill *et al.*, "On the opportunities and risks of foundation models," *arXiv preprint arXiv:2108.07258*, 2021.
- [18] D. L. K. Yamins and J. J. DiCarlo, "Using goal-driven deep learning models to understand sensory cortex," *Nature Neuroscience*, vol. 19, no. 3, pp. 356–365, 2016.
- [19] N. Kriegeskorte, "Deep neural networks: A new framework for modeling biological vision and brain information processing," *Annual Review of Vision Science*, vol. 1, no. 1, pp. 417–446, 2015.
- [20] M. Schrimpf, I. A. Blank, G. Tuckute, C. Kauf, E. A. Hosseini, N. Kanwisher, J. B. Tenenbaum, and E. Fedorenko, "The neural architecture of language: Integrative modeling converges on predictive processing," *Proceedings of the National Academy of Sciences*, vol. 118, no. 45, p. e2105646118, 2021.
- [21] D. C. Van Essen, S. M. Smith, D. M. Barch, T. E. J. Behrens, E. Yacoub, K. Ugurbil, and WU-Minn HCP Consortium, "The WU-Minn human connectome project: An overview," *NeuroImage*, vol. 80, pp. 62–79, 2013.
- [22] M. F. Glasser, S. N. Sotiropoulos, J. A. Wilson, T. S. Coalson, B. Fischl, J. L. Andersson, J. Xu, S. Jbabdi, M. Webster, J. R. Polimeni, D. C. Van Essen, M. Jenkinson, and WU-Minn HCP Consortium, "The minimal preprocessing pipelines for the human connectome project," *NeuroImage*, vol. 80, pp. 105–124, 2013.
- [23] E. C. Robinson, S. Jbabdi, M. F. Glasser, J. Andersson, G. C. Burgess, M. P. Harms, S. M. Smith, D. C. Van Essen, and M. Jenkinson, "MSM: A new flexible framework for multimodal surface matching," *NeuroImage*, vol. 100, pp. 414–426, 2014.
- [24] M. F. Glasser, T. S. Coalson, E. C. Robinson, C. D. Hacker, J. Harwell, E. Yacoub, K. Ugurbil, J. Andersson, C. F. Beckmann, M. Jenkinson, S. M. Smith, and D. C. Van Essen, "A multi-modal parcellation of human cerebral cortex," *Nature*, vol. 536, no. 7615, pp. 171–178, 2016.
- [25] J. L. Ji, M. Spronk, K. Kulkarni, G. Repovš, A. Anticevic, and M. W. Cole, "Mapping the human brain's cortical-subcortical functional network organization," *NeuroImage*, vol. 185, pp. 35–57, 2019.
- [26] B. T. T. Yeo, F. M. Krienen, J. Sepulcre, M. R. Sabuncu, D. Lashkari, M. Hollinshead, J. L. Roffman, J. W. Smoller, L. Zöllei, J. R. Polimeni, B. Fischl, H. Liu, and R. L. Buckner, "The organization of the human cerebral cortex estimated by intrinsic functional connectivity," *Journal of Neurophysiology*, vol. 106, no. 3, pp. 1125–1165, 2011.
- [27] D. Lakens, "Equivalence tests: A practical primer for t tests, correlations, and meta-analyses," *Social Psychological and Personality Science*, vol. 8, no. 4, pp. 355–362, 2017.
- [28] J. N. Rouder, P. L. Speckman, D. Sun, R. D. Morey, and G. Iverson, "Bayesian t tests for accepting and rejecting the null hypothesis," *Psychonomic Bulletin & Review*, vol. 16, no. 2, pp. 225–237, 2009.
- [29] V. D. Blondel, J.-L. Guillaume, R. Lambiotte, and E. Lefebvre, "Fast unfolding of communities in large networks," *Journal of Statistical Mechanics: Theory and Experiment*, vol. 2008, no. 10, p. P10008, 2008.
- [30] R. Guimerà and L. A. N. Amaral, "Functional cartography of complex metabolic networks," *Nature*, vol. 433, no. 7028, pp. 895–900, 2005.
- [31] J. Bai, S. Bai, Y. Chu, Z. Cui, K. Dang, X. Deng, Y. Fan, W. Ge, Y. Han, F. Huang *et al.*, "Qwen technical report," *arXiv preprint arXiv:2309.16609*, 2023.
- [32] D. Dai, C. Deng, C. Zhao, R. X. Xu, H. Gao, D. Chen, J. Li, W. Zeng, X. Yu, Y. Wu *et al.*, "Deepseekmoe: Towards ultimate expert specialization in mixture-of-experts language models," *arXiv preprint arXiv:2401.06066*, 2024.
- [33] G. Lai, Q. Xie, H. Liu, Y. Yang, and E. Hovy, "Race: Large-scale reading comprehension dataset from examinations," *arXiv preprint arXiv:1704.04683*, 2017.
- [34] S. Bowman, G. Angeli, C. Potts, and C. D. Manning, "A large annotated corpus for learning natural language inference," in *Proc. EMNLP*, 2015, pp. 632–642.
- [35] R. Socher, A. Perelygin, J. Wu, J. Chuang, C. D. Manning, A. Y. Ng, and C. Potts, "Recursive deep models for semantic compositionality over a sentiment treebank," in *Proc. EMNLP*, 2013, pp. 1631–1642.
- [36] K. Cobbe, V. Kosaraju, M. Bavarian, M. Chen, H. Jun, L. Kaiser, M. Plappert, J. Tworek, J. Hilton, R. Nakano *et al.*, "Training verifiers to solve math word problems," *arXiv preprint arXiv:2110.14168*, 2021.
- [37] N. Mantel, "The detection of disease clustering and a generalized regression approach," *Cancer Research*, vol. 27, no. 2, pp. 209–220, 1967.
- [38] V. A. Traag, L. Waltman, and N. J. van Eck, "From Louvain to Leiden: guaranteeing well-connected communities," *Scientific Reports*, vol. 9, p. 5233, 2019.